Horizon — research workflow (./run_all.sh)

quick = default · pilot = + small real runs · full = + every §11 experiment

1 Setup
scripts/setup.sh: uv · Python 3.11 · uv.lock · CaMeL (pinned) · Ollama + Qwen3-8B · Jupyter kernel

2 Data + data card
BIPIA · LLMail-Inject · HarmBench · JBB · deepset · Gandalf · Tensor Trust · XSTest · Alpaca

3 Fixed splits
fit → choose τ^*, π^*, $\hat{\lambda}$ · report → evaluate at fixed choices · paired design

4 Model serving
Ollama (GPU) + developer→system role proxy · colibri (report model)

5 Unit tests (§9)
77 tests · slide values only in tests/ (grep check)

6 Theory-mode figures
F1–F15 from params.yaml (SIMULATED inputs labelled)

7 Smoke test
quarantine + vote · policy · Z3 refusal · canary · guards

8 Experiments (level pilot / full)

CaMeL baseline
reference code, unmodified

Stage 1 runner
paired A[x, τ] · utility U_τ · adaptive attacker
fit: π^*, τ^* → report: evaluated

pi_star.json

Calibrate $\hat{\lambda}$
CRC on the calibration split

lambda_hat.json

AgentDojo ablations
none · Stage 1 · Stage 2 · Stage 1+2 · full
static, then adaptive attacks

Jailbreak runner
detection / FPR · PyRIT PAIR, TAP, Crescendo · garak

9 Figures
F1–F18 · PDF + PNG + CSV · formula + parameter footer

MLflow
config · seed · model · split · per-example outcomes

10 Notebook
Horizon.ipynb executed: every step with outputs

11 Model-written report
Qwen3.6-35B-A3B (colibri) analyses every figure, CSV, result tables from files · numbers checked → LaTeX → report.pdf